

String edit distance

Data Structures and Algorithms for Computational Linguistics III
(ISCL-BA-07)

Çağrı Çöltekin
ccoltekin@ifa.uni-tuebingen.de

University of Tübingen
Seminar für Sprachwissenschaft

Winter Semester 2020/21

version: 007020-0020-00-07

Hamming distance

a simple distance metric between two sequences

- The Hamming distance measures number of different symbols in the corresponding positions

h	y	g	i	e	n	e
h	i	g	i	e	n	e

0 + 1 + 0 + 0 + 0 + 0 = 1

h	y	g	i	e	n	e
h	i	y	g	e	i	n

0 + 1 + 1 + 1 + 0 + 1 + 1 = 5

- Very easy/efficient calculation
- But cannot handle sequences of different lengths (consider *hygme* – *hyggene*)

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 2 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

Edit distance

- In many applications, we want to know how similar (or different) two string are
 - Comparing two files (e.g., source code)
 - Comparing two DNA sequences
 - Spell checking
 - Approximate string matching
 - Determining similarity of two languages
 - Machine translation
- The solution is typically formulated as the (inverse) cost of obtaining one of the strings from the other through a number of *edit operations*
- Once we obtain the optimal edit operations, we may (depending on the edit operations) also be able to determine the optimal alignment between the strings

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 1 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

A family of edit distance problems

- The same overall idea applies to a number of well-known problems/solutions that differ in the type of operations allowed
 - Hamming distance*: only replacements
 - Longest common subsequence* (LCS) insertions and deletions
 - Levenshtein distance* insertions, deletions and substitutions
 - Levenshtein-Damerau distance* insertions, deletions and substitutions and transpositions (swap) of adjacent symbols
- Naive solutions to all (except Hamming distance) have exponential time complexity
- Polynomial-time solution can be obtained using *dynamic programming*

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 2 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

Longest common subsequence (LCS)

Problem definition

- A subsequence is an order-preserving (but not necessarily continuous) sequence of symbols from a string (a version of the sequence where zero or more elements are removed)
 - hgg*, *gn*, *yeie*, *heie* are subsequences of *hyggene*
- Note that a subsequence does not have to be a substring (substrings are continuous)
 - hgg*, *gicne*, *ene* are substrings of *hyggene*
- The longest common subsequence (LCS) of two strings is the longest string that is a subsequence of both strings
 - LCS(*hyggene*, *hyigene*) = *hyigene*
 - LCS(*hyigene*, *hyggene*) = *hyigene* / *hygene*
- LCS is exactly the problem solved by the UNIX *diff* utility
- It has wide-ranging applications from source-code comparison to bioinformatics (e.g., DNA sequencing)

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 3 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

LCS: a naive solution

- A simple solution is:
 - Enumerate all subsequences of the first string
 - Check if it is also a subsequence of the second string
- There are exponential number of subsequences of a string
 - the string *abc* has 8 subsequences:
 - abc*: nothing removed
 - ab*, *ac*, *bc*: individual elements are removed
 - a*, *b*, *c*: length-2 subsequences are removed
 - ε* (empty string): *abc* removed
 - For *abcd*, we have subsequences of *abc* once with, and once without *d*
 - Each additional symbol doubles the number of subsequences
- For strings of size *n* and *m*, the complexity of the brute-force algorithm is $O(2^n \cdot 2^m)$

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 4 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

LCS: recursive solution

demonstration

- Consider two strings *Xx*, *Yy* and their LCS *Zz* (*X*, *Y*, *Z* are possibly empty strings, *x*, *y*, *z* are characters)
- If *x = y*, then this character has to be part of the LCS, *x = y = z*, and *Z* must be the LCS of *X* and *Y*
- If *x ≠ y*, there are three cases
 - x ≠ y ≠ z*: *z* is also the LCS of *X* and *Y*
 - x = z*: *z* is also the LCS of *Xx* and *Y*
 - y = z*: *z* is also the LCS of *X* and *Yy*
- This leads to following recursive definition:

$$\text{LCS}(Xx, Yy) = \begin{cases} \text{LCS}(X, Y) & \text{if } x = y \\ \text{longer of } \text{LCS}(Xx, Y) \text{ and } \text{LCS}(X, Yy) & \text{otherwise} \end{cases}$$

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 5 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

LCS: divide-and-conquer



- Note the **repeated computation**

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 6 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

LCS: dynamic programming

general sketch

- For string indexes *i* and *j*, of strings *X* and *Y*, if we need $\text{LCS}(X_{1..i}, Y_{1..j})$, $\text{LCS}(X_{1..i-1}, Y_j)$, $\text{LCS}(X_i, Y_{j-1})$
- In the standard algorithm, we do not store the LCS, but the length of the LCS, $l_{i,j}$ for each *i, j*
- Once we fill in the matrix, the $l_{n,m}$ is the length of the LCS
- We can trace back and recover the LCS using the dynamic programming matrix

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 7 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

LCS with dynamic programming

demonstration

		0	1	2	3	4	5	6	7	8
		c	h	i	y	g	e	i	n	e
0	c	0	0	0	0	0	0	0	0	0
1	h	0	1	1	1	1	1	1	1	1
2	y	0	1	1	2	2	2	2	2	2
3	g	0	1	1	2	3	3	3	3	3
4	i	0	1	2	2	3	3	4	4	4
5	e	0	1	2	2	3	4	4	4	5
6	n	0	1	2	2	3	4	4	5	5
7	e	0	1	2	2	3	4	4	5	6

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 8 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

Complexity of filling the LCS matrix

```
1 = np.zeros(shape=(n+1,m+1))
for i in range(n):
    for j in range(m):
        if X[i] == Y[j]:
            l[i+1,j+1] = l[i,j] + 1
        else:
            l[i+1,j+1] = max(l[i+1,j], l[i,j+1])
```

- Two loops up to *n* and *m*, the time complexity is $O(nm)$
- Similarly, the space complexity is also $O(nm)$

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 9 / 19

Introduction/ motivation Longest common subsequence Levenshtein distance

Recovering the LCS from the matrix

Diagram showing the recovery of the LCS from the matrix. A blue path is traced back from the bottom-right cell (6,7) to the top-left cell (0,0), following the maximum values. The path is: (6,7) -> (5,7) -> (4,7) -> (3,7) -> (2,7) -> (1,7) -> (0,7) -> (0,6) -> (0,5) -> (0,4) -> (0,3) -> (0,2) -> (0,1) -> (0,0).

Ç. Çöltekin, SS / University of Tübingen

Winter Semester 2020/21 10 / 19

Transforming one string to another

- The table (back arrows) also gives a set of edit operations to transform one string to another
- For LCS, operations are:
 - copy (diagonal arrows in the demonstration)
 - insert (left arrows in the demonstration – assuming original string is the vertical one)
 - delete (up arrows in the demonstration – assuming original string is the vertical one)
- These also form an alignment between two strings
- Different set of edit operations recovered will yield the same LCS, but different alignments

LCS alignments

		0	1	2	3	4	5	6	7	8
		ε	h	i	y	g	e	i	n	e
0	ε	0	0	0	0	0	0	0	0	0
1	h	0	1	1	1	1	1	1	1	1
2	y	0	1	1	2	2	2	2	2	2
3	g	0	1	1	2	3	3	3	3	3
4	i	0	1	2	2	3	3	4	4	4
5	e	0	1	2	2	3	4	4	4	5
6	n	0	1	2	2	3	4	4	5	5
7	e	0	1	2	2	3	4	4	5	6

Alignments:

h-y-g-i-ne
 ciccicicc
 hiygei-ne
 h-y-g-e-ne
 ciccicicc
 hiyg-ei-ne

LCS – some remarks

- We formulated the algorithm as optimizing the LCS
- Alternatively, we can consider costs associated with each operation:
 - copy = 0
 - delete = 1
 - insert = 1
- This is the typical application of LCS, as in diff
- In some applications we may want to have different costs for delete and insert (e.g., mapping lemmas to inflected forms of words)
- Similarly, we may want to assign different costs for different characters (e.g., higher cost to delete consonants in historical linguistics)

Levenshtein distance

definition

- Levenshtein difference between two strings is the total cost of *insertions*, *deletions* and *substitutions*
- With cost of 1 for all operations

$$\text{lev}(Xx, Yy) = \begin{cases} \text{len}(X) & \text{if } \text{len}(y) = 0 \\ \text{len}(Y) & \text{if } \text{len}(x) = 0 \\ \text{lev}(X, Y) & \text{if } x = y \\ 1 + \min \begin{cases} \text{lev}(X, Yy) \\ \text{lev}(Xx, Y) \\ \text{lev}(X, Y) \end{cases} & \text{otherwise} \end{cases}$$

- Naive recursion (as defined above), again, is intractable
- But, the same dynamic programming method works

Levenshtein distance

demonstration

		0	1	2	3	4	5	6	7	8
		ε	h	i	y	g	e	i	n	e
0	ε	0	1	2	3	4	5	6	7	8
1	h	1	0	1	2	3	4	5	6	7
2	y	2	1	1	1	2	3	4	5	6
3	g	3	2	2	2	1	2	3	4	5
4	i	4	3	2	3	2	2	2	3	4
5	e	5	4	3	3	3	2	3	3	4
6	n	6	5	4	4	4	3	3	3	4
7	e	7	6	5	5	5	4	4	4	3

Levenshtein distance

edits and alignments

		0	1	2	3	4	5	6	7	8
		ε	h	i	y	g	e	i	n	e
0	ε	0	1	2	3	4	5	6	7	8
1	h	1	0	1	2	3	4	5	6	7
2	y	2	1	1	1	2	3	4	5	6
3	g	3	2	2	2	1	2	3	4	5
4	i	4	3	2	3	2	2	2	3	4
5	e	5	4	3	3	3	2	3	3	4
6	n	6	5	4	4	4	3	3	3	4
7	e	7	6	5	5	5	4	4	4	3

Edit distance: extensions and variations

- Another possible operation we did not cover is *swap* (or transpose), which is useful for applications like spell checking
- In some applications (e.g., machine translation, OCR correction) we may want to have one-to-many or many-to-one alignments
- Additional requirements often introduce additional complexity
- It is sometimes useful to learn costs from data

Summary

- Edit distance is an important problem in many fields including computational linguistics
- A number of related problems can be efficiently solved by dynamic programming
- Edit distance is also important for approximate string matching and alignment
- Reading suggestion: Goodrich, Tamassia, and Goldwasser (2013, chapter 13), Jurafsky and Martin (2009, section 3.11, or 2.5 in online draft)

Next:

- Algorithms on strings: tries
- Reading: Goodrich, Tamassia, and Goldwasser (2013, chapter 13),

Acknowledgments, credits, references

- Goodrich, Michael T., Roberto Tamassia, and Michael H. Goldwasser (2013). *Data Structures and Algorithms in Python*. John Wiley & Sons, Incorporated. [9781118476734](https://doi.org/10.1111/9781118476734).
- Jurafsky, Daniel and James H. Martin (2009). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*, second edition. Pearson Prentice Hall. [978-0-13-504196-3](https://doi.org/10.1111/9780135041963).

